

— EDITION 18

I came up in payment ops. Here's what I'm building.

Ten years in payment operations taught me reconciliation, idempotency and replay protection. That turned out to be exactly what running production AI takes.

A 7-part breakdown

SWIPE >>>>

0101 / 07

| From payment ops to production AI.

I spent a decade in payment operations. Now I run a fleet of AI agents on real workflows. The skills carried over more than I expected.

0202 / 07

What payment ops actually teaches you.

Batch windows. Idempotency. Replay protection. Reconciliation. The habit of assuming any job will get re-run and any file can fail halfway.

0303 / 07

| The first system I shipped.

A content engine. Around two dozen scheduled jobs that draft, queue, verify and publish across platforms, then pull the numbers back in. No team, just the agents and the wiring.

0404 / 07

| What broke first.

Context-window exhaustion. Jobs force-loaded big reference files, overflowed mid-run, re-read the same files and aborted. A classic resource leak in a new place.

0505 / 07

| How I fixed it.

Read slim. Rotate logs. Edit data with scripts, not the model. Post, then record, with a duplicate guard. The same controls I used on payment files.

0606 / 07

| Why it is the same job.

Running agents in production is operations work. Scheduling, finality, reconciliation, and a hard respect for the limits of the system. The model is the easy part.

0707 / 07

| Save this.

Payment ops to AI ops is not a leap. It is the same discipline with a new constraint budget.

— THE TAKEAWAY

Future-proof yet?

Comment what you're learning this quarter — let's compare notes. Repost to share it, and follow for daily agentic payments breakdowns.

◆ Joseph Bankole