

— EDITION 31

I debugged a production AI failure. The bug was the context window.

Two of my automated agents started failing the same way. The root cause was something I do not see discussed enough: context-window exhaustion as a production failure mode.

A 7-part breakdown

SWIPE >>>>

0101 / 07

| The setup.

I run a small fleet of automated AI agents: content pipelines, analytics jobs, scheduled workflows. Last week two of them started failing the same way.

0202 / 07

| The failure mode.

Context-window exhaustion. Not a model problem, a resource problem. The kind I know from payment operations.

0303 / 07

| Job one, what happened.

A daily job force-loaded around 141KB of reference files before doing any work. Add tool outputs and the working memory overflowed mid-run, compacted, re-read the same files, and overflowed again. Three loops, then abort.

0404 / 07

Job two, what happened.

A second job loaded a 25KB CSV, rewrote the whole file back, and crashed between posting and recording that it had posted. One rerun away from publishing a duplicate.

0505 / 07

The fixes were classic ops.

Read slim: a 63KB style guide became a 3KB cheatsheet. Rotate logs: anything over about 400 lines archives itself and jobs read the tail.

0606 / 07

| The other two fixes.

Edit data with scripts, not the model, so the file never enters memory. Post, then record immediately, with a duplicate guard that checks before acting.

0707 / 07

| Save this.

Ten years in payment operations taught me these patterns: batch windows, replay protection, reconciliation. Running agents in production is the same discipline with a new constraint budget.

— THE TAKEAWAY

Found this **useful?**

Repost to share it with your network, and follow for weekly breakdowns of the agentic payments stack.

◆ Joseph Bankole