

— EDITION 32

The hard part of running AI agents is never the model.

I run a content operation that is almost entirely automated. The writing was the easy bit. Everything wrapped around the model is where the work lives.

A 7-part breakdown

SWIPE >>>>

0101 / 07

| The setup.

I run a content operation with no team.
Scheduled AI agents generate posts, queue
them across platforms, pull analytics back in,
and adjust.

0202 / 07

The assumption everyone makes.

People assume the hard part is the writing. It isn't. The models write fine. The hard part is everything around them.

0303 / 07

| Orchestration.

Two dozen scheduled jobs that have to not step on each other. One builds tomorrow's slate at night, another verifies and ships it in the morning, others pull stats. None of them duplicating work.

0404 / 07

| Idempotency.

A scheduled job will get re-run. If "did I already do this?" isn't built in, you publish twice. You build the check before you build the feature.

0505 / 07

| Evals as a log, not a vibe.

I keep a performance log per post, so the system breeds toward what actually works instead of what feels good. Without it you're tuning on a hunch.

0606 / 07

| What nearly killed it.

Not the AI. Treating reference files and logs as free to load. They aren't. They eat the context budget until a run can't think. Fixing that was pure operations discipline.

0707 / 07

| Save this.

Running a fleet of agents is the same job I came up in: scheduling, reconciliation, replay protection, and a hard respect for the limits of the system.

— THE TAKEAWAY

Found this **useful?**

Repost to share it with your network, and follow for weekly breakdowns of the agentic payments stack.

◆ Joseph Bankole