

— EDITION 40

A scheduled agent that writes 15 posts a week, then checks its own work

I built a content agent that writes fifteen posts a week against one fixed brief, then grades every one before it ships.

A 7-part breakdown

SWIPE >>>>

0101 / 07

| One fixed brief.

Fifteen posts a week, written against a single document instead of fifteen separate prompts. Consistency comes from the brief, not from me reviewing every draft by hand.

0202 / 07

| A word count per post.

Short posts stay short and long ones stay long, because a ceiling is set before the writing starts. No draft drifts to whatever length felt right in the moment.

0303 / 07

A type mix across five formats.

Practitioner story, debugging note, deployment update and two others, tracked across the week so one format cannot quietly crowd out the rest.

0404 / 07

No repeated hook structures.

The agent checks the last two weeks of hooks before writing a new one. If the opening line pattern has already run, it has to find a different way in.

0505 / 07

| The eight-point gate.

Every finished post is checked against eight rules before it is allowed to queue. A fail sends it back to be regenerated, not lightly edited.

0606 / 07

The rule that mattered most.

No invented numbers, ever. If a stat cannot be verified, the post ships without one rather than with a plausible-sounding guess.

0707 / 07

| Sourcing before writing.

That one constraint moved sourcing to the front of the pipeline. You cannot verify a number you look for after the sentence is already written.

— THE TAKEAWAY

That is the brief, the mix and the **gate.**

What is the one rule in your pipeline that quietly shapes everything downstream of it? Follow Joseph Bankole for more from inside a working AI content operation.

◆ Joseph Bankole