◆ **JOSEPH** BANKOLE

FIELD NOTES

EDITION 43

Nine scheduled agents, and three were doing the same job

Nine agents. Three of them re-checked the same source data every morning, separately. So I deleted four.

A **7**-part breakdown

SWIPE >>>>

0101 / 07

Nine scheduled agents.

The content operation ran as nine separate scheduled jobs, each with its own trigger time and its own task.

0202 / 07

Three did the same job.

Every morning, three of them independently re-verified the same source data before any copy got written.

0303 / 07

Same facts, three times.

Three separate verification passes over identical results, and three separate token bills for one answer.

0404 / 07

One shared intel file.

The fix was a single morning desk that verifies once and writes the results to a shared file the rest of the pipeline reads from.

0505 / 07

Nine sessions became five.

Merging the morning cluster into one desk cut four redundant jobs straight out of the schedule.

0606 / 07

Cadence did not change.

Posting frequency stayed identical. Every verify gate and every humanizer pass that mattered stayed in place, just run once instead of three times.

0707 / 07

The real cost of a fleet.

It is rarely the model. It is the same work done in parallel by jobs that have no idea the others exist.

THE TAKEAWAY

Fewer agents, same output.

How do you stop agents on separate schedules from redoing each other's work? Follow Joseph Bankole for more from running an AI content operation.

◆ Joseph Bankole

FOLLOW »»