

— EDITION 44

Order the pipeline for failure, not for logic

*Order your pipeline for failure, not for logic.
Here is what it cost me to learn that.*

A 7-part breakdown

SWIPE >>>>

0101 / 07

| One job, two tasks.

A single scheduled agent ships the day's post slate to the queue and also renders a heavy motion clip through an external generation API.

0202 / 07

| The render is the flaky part.

It stalls without warning, roughly once a week, with no pattern to when.

0303 / 07

| It sat in front of the posts.

In the old order, the render ran first.
Everything else in the job waited behind it.

0404 / 07

A stall held the whole queue.

When the render hung, the day's post slate never shipped either, because it was queued behind a step that had not finished.

0505 / 07

| I reordered the job.

Slate and threads go out first now. The render, the least reliable step, runs last.

0606 / 07

A stall now costs one reel.

If the render hangs today, it delays a single video. It no longer touches the posts that were ready hours earlier.

0707 / 07

| Rank by reliability, not logic.

Put the step most likely to break where its failure costs the least, even if that is not the order the work makes sense in your head.

— THE TAKEAWAY

Order for failure, not for **flow.**

Does the flakiest step in your pipeline run early or late, and what happens to everything behind it when it hangs? Follow Joseph Bankole for more from building AI systems in public.

◆ Joseph Bankole