

FIELD NOTES

Running a Fleet of AI Agents.

The model is the easy part. Everything around it is the job.

Swipe >>>

THE STACK · ORCHESTRATION

Stop the jobs colliding

Two dozen scheduled jobs. One drafts at night, another verifies and ships at dawn, none of them doing the same work twice.

THE STACK · BUDGET

Treat context as finite

Every job runs inside a context window and a credit budget. Treat them as free and the whole thing falls over. It did.

THE STACK · IDEMPOTENCY

Assume it re-runs

A scheduled job will get re-run. If 'did I already do this?' is not built in, you publish twice.

THE STACK · EVALS

Keep score, not vibes

A performance log per post, so the system breeds toward what actually works instead of what felt clever.

WHAT NEARLY KILLED IT

The context leak

Loading reference files and logs as if they were free. They eat the budget until a run cannot think. The fix was operations discipline.

THE TAKEAWAY

The model is the easy part.

If you orchestrate agents on real work,
what is harder than the model itself?

◆ Follow Joseph Bankole for field notes